

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error
import warnings
```


```
In [2]: data= pd.read_csv("C:\\Users\\AMIT\\Downloads\\data (2).csv")
```

```
In [3]: data.head()
```

Out[3]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
0	842302	M	17.99	10.38	122.80	1001.0	0.11840
1	842517	M	20.57	17.77	132.90	1326.0	0.08474
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960
3	84348301	M	11.42	20.38	77.58	386.1	0.14250
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030

5 rows × 33 columns



In [4]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    569 non-null    int64
1   diagnosis                            569 non-null    object
2   radius_mean                          569 non-null    float64
3   texture_mean                         569 non-null    float64
4   perimeter_mean                       569 non-null    float64
5   area_mean                           569 non-null    float64
6   smoothness_mean                      569 non-null    float64
7   compactness_mean                     569 non-null    float64
8   concavity_mean                       569 non-null    float64
9   concave points_mean                  569 non-null    float64
10  symmetry_mean                        569 non-null    float64
11  fractal_dimension_mean               569 non-null    float64
12  radius_se                            569 non-null    float64
13  texture_se                           569 non-null    float64
14  perimeter_se                         569 non-null    float64
15  area_se                              569 non-null    float64
16  smoothness_se                        569 non-null    float64
17  compactness_se                       569 non-null    float64
18  concavity_se                         569 non-null    float64
19  concave points_se                    569 non-null    float64
20  symmetry_se                          569 non-null    float64
21  fractal_dimension_se                 569 non-null    float64
22  radius_worst                         569 non-null    float64
23  texture_worst                        569 non-null    float64
24  perimeter_worst                      569 non-null    float64
25  area_worst                           569 non-null    float64
26  smoothness_worst                     569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                 569 non-null    float64
30  symmetry_worst                       569 non-null    float64
31  fractal_dimension_worst              569 non-null    float64
32  Unnamed: 32                          0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

In [5]: data.isna().sum()

```
Out[5]: id                                0
diagnosis                                0
radius_mean                             0
texture_mean                             0
perimeter_mean                           0
area_mean                                0
smoothness_mean                           0
compactness_mean                           0
concavity_mean                             0
concave points_mean                         0
symmetry_mean                             0
fractal_dimension_mean                     0
radius_se                                 0
texture_se                                 0
perimeter_se                               0
area_se                                   0
smoothness_se                             0
compactness_se                             0
concavity_se                               0
concave points_se                         0
symmetry_se                               0
fractal_dimension_se                       0
radius_worst                              0
texture_worst                             0
perimeter_worst                           0
area_worst                                0
smoothness_worst                           0
compactness_worst                           0
concavity_worst                             0
concave points_worst                       0
symmetry_worst                             0
fractal_dimension_worst                     0
Unnamed: 32                               569
dtype: int64
```

In [6]: data.describe()

Out[6]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
<b>count</b>	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000
<b>mean</b>	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360
<b>std</b>	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014064
<b>min</b>	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630
<b>25%</b>	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370
<b>50%</b>	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870
<b>75%</b>	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300
<b>max</b>	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400

8 rows × 32 columns


In [7]: df=data.drop(columns='Unnamed: 32')

In [8]: `df.shape`

Out[8]: (569, 32)

In [9]: `sns.countplot(x='diagnosis' , hue='diagnosis', data=df)`

Out[9]: <AxesSubplot:xlabel='diagnosis', ylabel='count'>

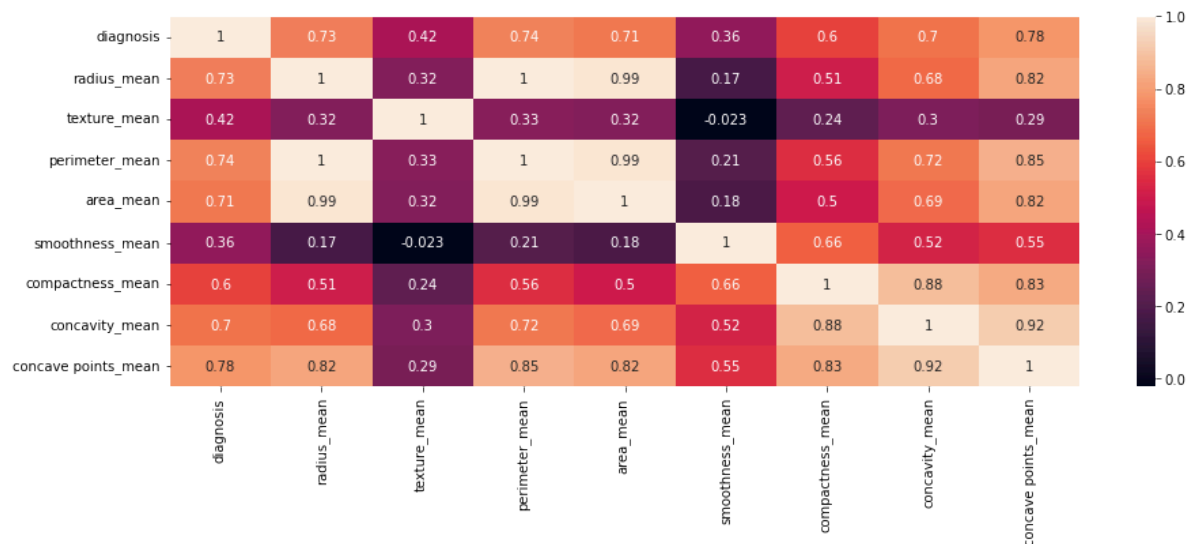


In [10]: `lb= LabelEncoder()`

In [11]: `df.iloc[:, 1]=lb.fit_transform(df.iloc[:, 1].values)`

In [12]: `plt.figure(figsize=(15,5))`  
`sns.heatmap(df.iloc[:, 1:10].corr(), annot=True)`

Out[12]: <AxesSubplot:>



In [13]: `X= df.iloc[:, 2:32].values`  
`y=df.iloc[:, 1].values`

In [14]: `X_train, X_test,y_train, y_test=train_test_split(X,y,test_size=0.2,random_state=0)`

```
In [15]: st=StandardScaler()
X_train=st.fit_transform(X_train)
X_test=st.fit_transform(X_test)
```

```
In [16]: data.drop(['Unnamed: 32'], axis=1, inplace=True)
y = data['diagnosis'].map({'M': 1, 'B': 0})
X = data.drop(['id', 'diagnosis'], axis =1)
```

```
In [17]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Initialize Random Forest classifier
rf = RandomForestClassifier()

# Train the Random Forest classifier
rf.fit(X_train, y_train)

# Check the accuracy on the training set
train_accuracy = rf.score(X_train, y_train)

# Make predictions on the test set and calculate accuracy
test_predictions = rf.predict(X_test)
test_accuracy = accuracy_score(y_test, test_predictions)

# Print classification report
print(classification_report(y_test, test_predictions))

print("Training Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	67
1	0.98	0.96	0.97	47
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

Training Accuracy: 1.0


Test Accuracy: 0.9736842105263158

```
In [18]: from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, test_predictions)
cm
```


```
Out[18]: array([[66, 1],
                [ 2, 45]], dtype=int64)
```

```
In [19]: sns.heatmap(cm, annot=True)  
plt.show()
```




```
In [20]: from sklearn.metrics import confusion_matrix  
import matplotlib.pyplot as plt  
# Logistic Regression  
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)  
y_pred_logreg = logreg.predict(X_test)  
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)  
print("Logistic Regression Accuracy:", accuracy_logreg)  
cm_logreg = confusion_matrix(y_test, y_pred_logreg)  
sns.heatmap(cm_logreg, annot=True)  
plt.show()
```

Logistic Regression Accuracy: 0.956140350877193



```
In [21]: # K-Nearest Neighbors
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print("K-Nearest Neighbors Accuracy:", accuracy_knn)
cm_knn = confusion_matrix(y_test, y_pred_knn)
sns.heatmap(cm_knn, annot=True)
plt.show()
```

K-Nearest Neighbors Accuracy: 0.9649122807017544




```
In [22]: # Linear Regression
linear_regression = LinearRegression()
linear_regression.fit(X_train, y_train)
y_pred_linear_regression = linear_regression.predict(X_test)
mse_linear_regression = mean_squared_error(y_test, y_pred_linear_regression)
print("Linear Regression MSE:", mse_linear_regression)
```

Linear Regression MSE: 0.068366677209251

```
In [23]: # Support Vector Machine
from sklearn.svm import SVC
svm = SVC()
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print("Support Vector Machine Accuracy:", accuracy_svm)
cm_svm = confusion_matrix(y_test, y_pred_svm)
sns.heatmap(cm_svm, annot=True)
plt.show()
```

Support Vector Machine Accuracy: 0.9824561403508771






```
In [24]: import matplotlib.pyplot as plt
# Create a bar chart with different colors for each bar
colors = ["red", "green", "blue", "orange", "purple"]
labels = ["Logistic Regression", "K-Nearest Neighbors", "Random Forest", "Linear",
          "Support Vector Machine"]
accuracies = [accuracy_logreg, accuracy_knn, test_accuracy, mse_linear_regression, test_svm_accuracy]

plt.figure(figsize=(10, 6))
plt.bar(labels, accuracies, color=colors)

# Add title and axis labels
plt.title("Model Accuracies", fontsize=16)
plt.xlabel("Model", fontsize=14)
plt.ylabel("Accuracy", fontsize=14)

# Add grid and show plot
plt.grid(True)
plt.show()
```



```
In [26]: # Take user input for new data
new_data = []
print("Enter the values for each feature:")

# Prompt the user for each feature value
for col in X.columns:
    val = input(f"Enter {col}: ")
    new_data.append(float(val))

# Convert user input into numpy array
new_data = np.array(new_data).reshape(1, -1)

# Standardize the new data using the same scaler
new_data = st.transform(new_data)

# Predict using the trained Random Forest classifier
prediction = rf.predict(new_data)

# Display the prediction
if prediction[0] == 1:
    print("The diagnosis is Malignant (M)")
else:
    print("The diagnosis is Benign (B)")
```

Enter the values for each feature:

Enter radius\_mean: 567  
Enter texture\_mean: 67.8  
Enter perimeter\_mean: 76  
Enter area\_mean: 45  
Enter smoothness\_mean: 5656  
Enter compactness\_mean: 57  
Enter concavity\_mean: 788  
Enter concave points\_mean: 67  
Enter symmetry\_mean: 566  
Enter fractal\_dimension\_mean: 46  
Enter radius\_se: 45  
Enter texture\_se: 56  
Enter perimeter\_se: 90  
Enter area\_se: 89.6  
Enter smoothness\_se: 0.86  
Enter compactness\_se: 676  
Enter concavity\_se: 67  
Enter concave points\_se: 432  
Enter symmetry\_se: 35.9  
Enter fractal\_dimension\_se: 87.0  
Enter radius\_worst: 7.8  
Enter texture\_worst: 7.09  
Enter perimeter\_worst: 0.98  
Enter area\_worst: 87  
Enter smoothness\_worst: 6  
Enter compactness\_worst: 6  
Enter concavity\_worst: 7  
Enter concave points\_worst: 65  
Enter symmetry\_worst: 557  
Enter fractal\_dimension\_worst: 6768  
The diagnosis is Benign (B)

